

LECTURE 14

INTRODUCTION TO CONCURRENCY

SUBJECTS

Concepts of concurrency

Subprogram-level concurrency

Semaphores

Deadlocks

CONCEPT OF CONCURRENCY

Concurrency can be achieved at different levels:

- Instruction level
- Unit level
- Program level

Concurrent execution of program units can be performed:

- Physically on separate processors,
- or Logically in some time-sliced fashion on a single processor computer system

WHY STUDY CONCURRENCY?

Computing systems solve real world problems, where things happen at the same time:

- Railway Networks (lots of trains running, but that have to synchronize on shared sections of track)
- Operating System, with lots of processes running concurrently
- Web servers

Multiple processor or multi-core computers are now being widely used

- This creates the need for software to make effective use of that hardware capability

SUBPROGRAM-LEVEL CONCURRENCY

A **task is a unit of program that can be executed concurrently with other units of the same program**

- Each task in a program can provide one thread of control

A task can communicate with other tasks through:

- Shared nonlocal variables
- Message passing
- Parameters

SYNCHRONIZATION

A mechanism to control the order in which tasks execute

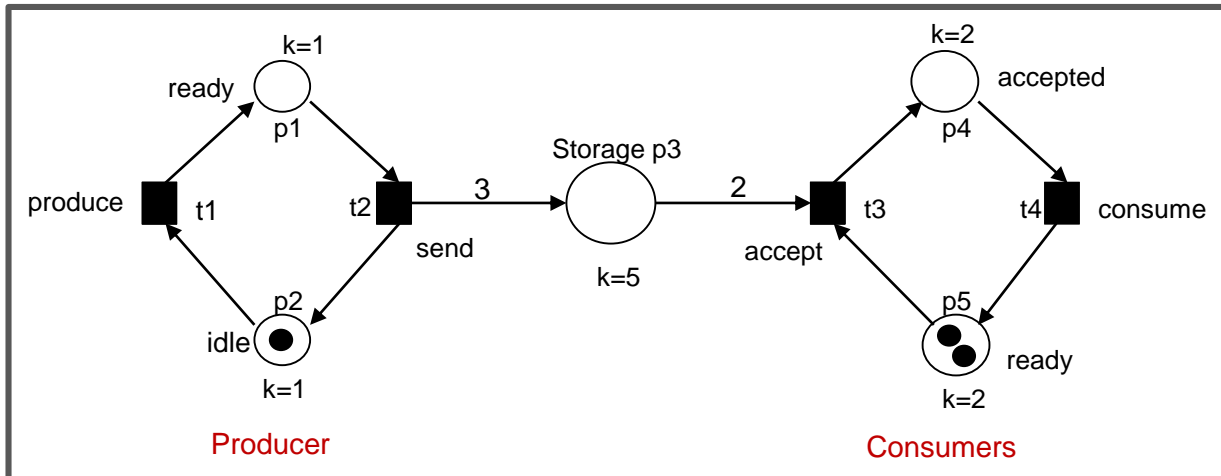
Cooperation synchronization is required between task A and task B when:

- **Task A** must wait for **task B** to complete some specific activity before **task A** can continue execution
- **Recall** the producer-consumer petri net problem


Competition synchronization is required between two tasks when:

- Both require the use of some resource that cannot be simultaneously used

SYNCHRONIZATION (PETRI NETS)

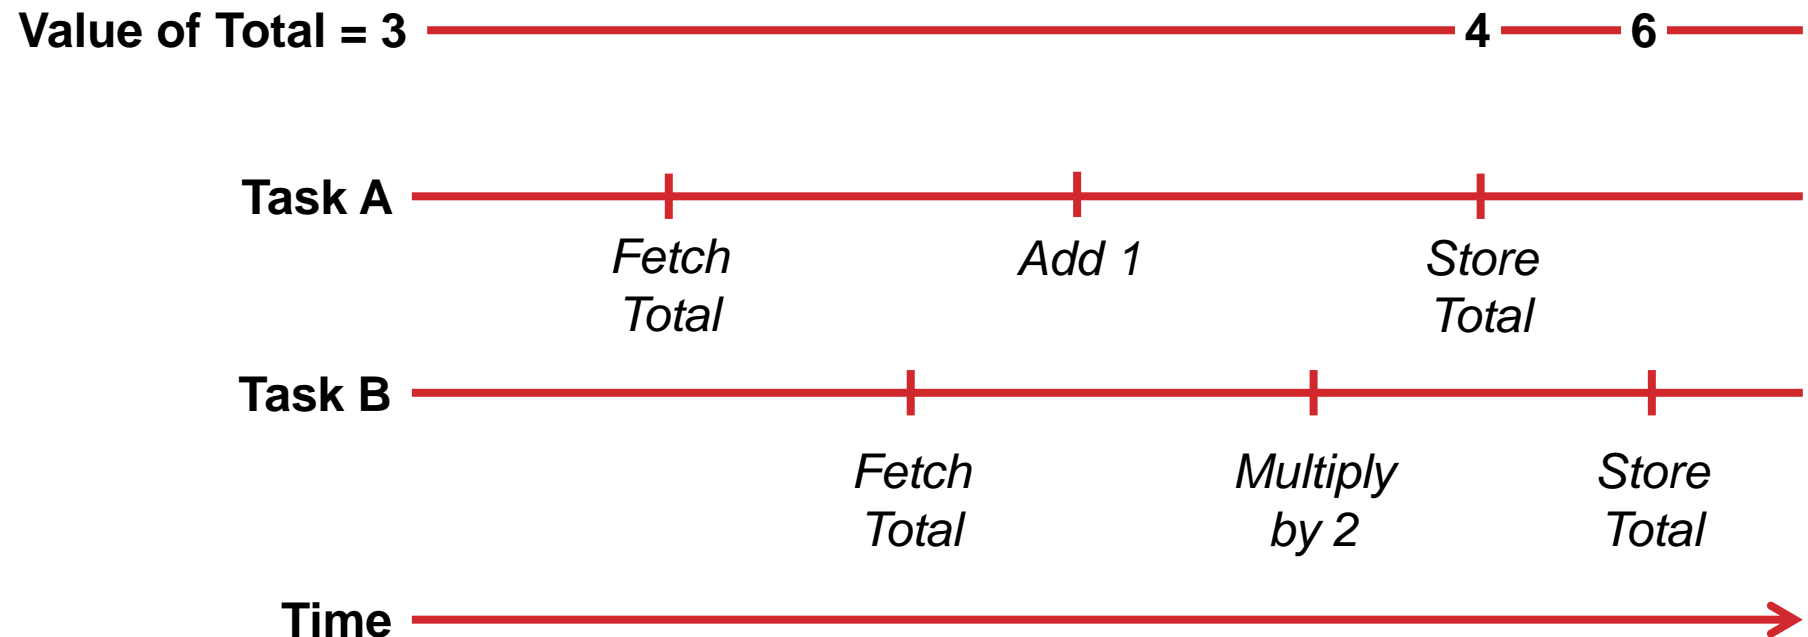


**Cooperation
Synchronization**



**Competition
Synchronization**

THE NEED FOR COMPETITION SYNCHRONIZATION



CRITICAL SECTION

A segment of code, in which the thread may be:

- Changing common variables,
- Updating a table,
- Writing to a file,
- Or updating any shared resource

The execution of critical sections by the threads is mutually exclusive in time

TASK (THREAD) STATES

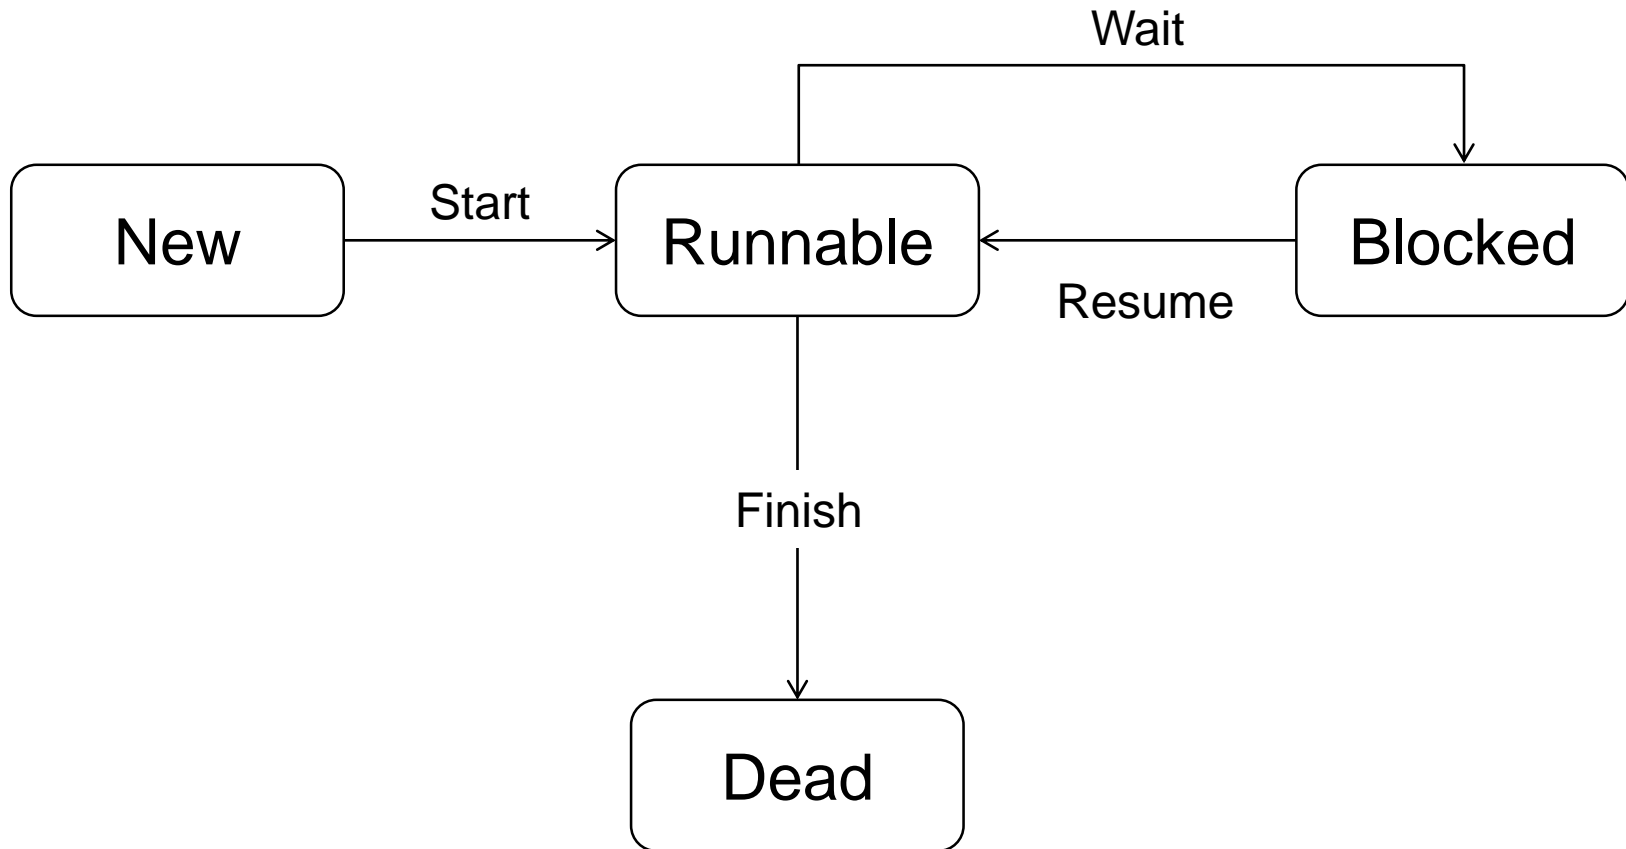
New: it has been created, but has not yet begun its execution

Runnable or **ready:** it is currently running or is ready to run

Blocked: it has been running, but its execution was interrupted by one of several different events

Dead: no longer active in any sense

TASK (THREAD) STATES



SEMAPHORES

Semaphore is a technique used to control access to a common resource for multiple tasks

- It is an object consisting of an integer (*counter*) and a queue that stores task descriptors

A task that requires access to a critical section needs to “acquire” the semaphore

Classically: a system of sending messages by holding the arms of two flags in certain positions according to an alphabetic code

SEMAPHORES

Two operations are always associated with Semaphores

Operation “**P**” is used to get the semaphore, and “**V**” to release it

- **P** (proberen, meaning to test and decrement the integer)
- **V** (verhogen, meaning to increment) operation

Alternatively, these operations are called: **wait** and **release**

SEMAPHORES

Operating systems often distinguish between counting and binary semaphores

The value of the counter of a **counting semaphore** can range over an unrestricted domain.

The value the counter of a **binary semaphore** can range only between 0 and 1

BINARY SEMAPHORES

The general strategy for using a binary semaphore to control access to a critical section is as follows:

```
Semaphore aSemaphore;
```

```
wait(aSemaphore);
```

```
Critical section();
```

```
release(aSemaphore);
```

BINARY SEMAPHORE - WAIT

wait(binarySemaphore) :

```
if binarySemaphore's counter == 1 then  
    set binarySemaphore's counter = 0  
else  
    Set the task's state to blocked  
    Put the caller in binarySemaphore's queue  
end
```


BINARY SEMAPHORE - RELEASE

release (binarySemaphore) :

```
if binarySemaphore's queue is empty then
    set binarySemaphore's counter = 1
else
    set the state of the task at the queue head to runnable
    Perform a dequeue operation
end
```

COUNTING SEMAPHORE - WAIT

```
wait(countingSemaphore) :
```

```
if countingSemaphore's counter > 0 then  
    Decrement countingSemaphore's counter  
else  
    Set the task's state to blocked  
    Put the caller in countingSemaphore's queue  
end
```

COUNTING SEMAPHORE - RELEASE

release(countingSemaphore) :

```
if countingSemaphore's queue is empty then  
    increment countingSemaphore's counter  
else  
    set the state of the task at the queue head to ready  
    Perform a dequeue operation  
end
```

PRODUCER AND CONSUMER

```
semaphore fullspots, emptyspots;
fullspots.count := 0;
emptyspots.count := BUFLen;
task producer;
  loop
    -- produce VALUE --
    wait(emptyspots);    { wait for a space }
    DEPOSIT(VALUE);
    release(fullspots);  { increase filled spaces }
  end loop;
end producer;

task consumer;
  loop
    wait(fullspots);     { make sure it is not empty }
    FETCH(VALUE);
    release(emptyspots); { increase empty spaces }
    -- consume VALUE --
  end loop
end consumer;
```

```
wait(countingSemaphore) :
```

```
if countingSemaphore's counter > 0 then
  Decrement countingSemaphore's counter
else
  Set the task's state to blocked
  Put caller in countingSemaphore's queue
end
```

```
release(countingSemaphore) :
```

```
if countingSemaphore's queue is empty then
  increment countingSemaphore's counter
else
  set state of task at queue head to ready
  Perform a dequeue operation
end
```

ADDING COMPETITION SYNCHRONIZATION

```
semaphore access, fullspots, emptyspots;
access.count := 1;
fullspots.count := 0;
emptyspots.count := BUFLen;

task producer;
  loop
    -- produce VALUE --
    wait(emptyspots);      { wait for a space }
    wait(access);          { wait for access }
    DEPOSIT(VALUE);
    release(access);       { relinquish access }
    release(fullspots);    { increase filled spaces }
  end loop;
end producer;
```


COMPETITION SYNCHRONIZATION (II)

```
task consumer;
  loop
    wait(fullspots);           { make sure it is not empty }
    wait(access);              { wait for access }
    FETCH(VALUE);
    release(access);           { relinquish access }
    release(emptyspots);       { increase empty spaces }
    -- consume VALUE --
  end loop
end consumer;
```

DEADLOCKS


A law passed by the Kansas legislature early in 20th century:

“..... When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone.”



DEADLOCK

Classic case of traffic deadlock. This is a gridlock where not vehicle can move forward to clear the traffic jam



CONDITIONS FOR DEADLOCK

Mutual exclusion: the act of allowing only one process to have access to a dedicated resource

Hold-and-wait: there must be a process holding at least one resource and waiting to acquire additional ones that are currently being held by other processes

No preemption: the lack of temporary reallocation of resources. Resources can be released only voluntarily

Circular waiting: each process involved in the impasse is waiting for another to voluntarily release the resource

DEADLOCK EXAMPLE

BinarySemaphore s1, s2;

Task A:

```
wait(s1)
    // access resource 1
wait (s2)
    // access resource 2
release(s2)
release(s1)
end task
```

Task B:

```
wait(s2)
    // access resource 2
wait(s1)
    // access resource 1
release(s1)
release(s2)
end task
```

STRATEGY FOR HANDLING DEADLOCKS

Prevention: eliminate one of the necessary conditions

Avoidance: avoid if the system knows ahead of time the sequence of resource requests associated with each active processes

Detection: detect by building directed resource graphs and looking for circles

Recovery: once detected, it must be untangled and the system returned to normal as quickly as possible

- Process termination
- Resource preemption

THANK YOU!

QUESTIONS?