

# LECTURE 17

## INTER-PROCESS COMMUNICATION AND PROCESS SCHEDULING

# SUBJECTS

**Unicast vs Multicast**

**Message Passing (Synchronous vs Asynchronous)**

**CPU scheduling**

- First Come First Served
- Round Robin

# IPC – UNICAST AND MULTICAST

**Distributed computing involves two or more processes engaging in IPC**

- They use a pre-agreed upon protocol

**A process may act as a sender at some point and a receiver at another**


**Unicast: communication between two processes**

- E.g., Socket communication


**Multicast: communication between one process and a group of processes**

- E.g., Publish/Subscribe Message model

# UNICAST VS. MULTICAST



**unicast**



**multicast**

# MESSAGE PASSING

**Message passing means that one process sends a message to another process and then continues its local processing**

- The message may take some time to get to the other process

**The message may be stored in the input queue of the destination process**

- If the latter is not immediately ready to receive the message

**Two types of message passing:**

- Asynchronous message passing
- Synchronous message passing

# ASYNCHRONOUS MESSAGE PASSING

## Blocking send and receive operations:

- A receiver will be blocked if it arrives at the point where it may receive messages and no message is waiting.
- A sender may get blocked if there is no room in the message queue between the sender and the receiver
  - However, in many cases, one assumes arbitrary long queues, which means that the sender will **almost** never be blocked

# ASYNCHRONOUS MESSAGE PASSING

## Non-blocking send and receive operations:

- Send and receive operations always return immediately
  - They return a status value which could indicate that no message has arrived at the receiver
- The receiver may test whether a message is waiting and possibly do some other processing
  - It may optionally be notified by the system when a message is received

# SYNCHRONOUS MESSAGE PASSING

**One assumes that sending and receiving takes place at the same time**

- There is often no need for an intermediate buffer

**This is also called rendezvous and implies closer synchronization:**

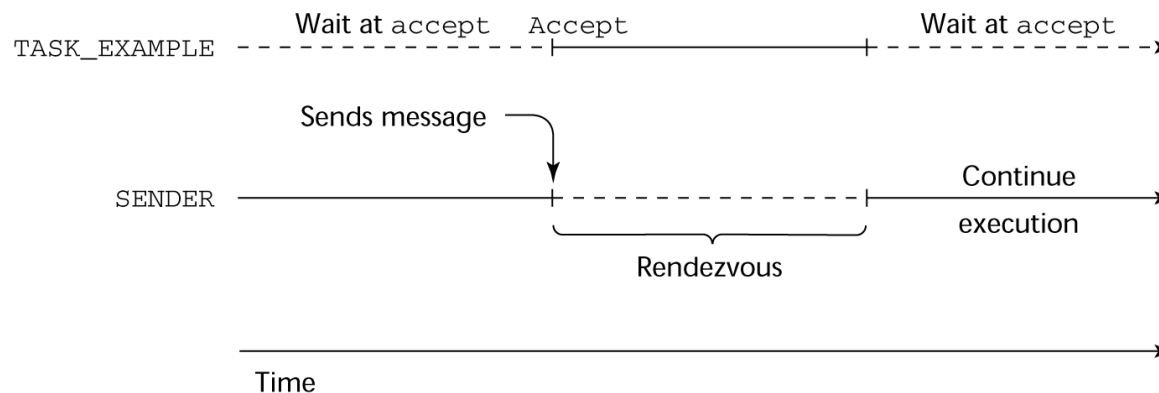
- The combined send-and-receive operation can only occur if both parties are ready to do their part.

**The sending process may have to wait for the receiving process, or the receiving process may have to wait for the sending one**

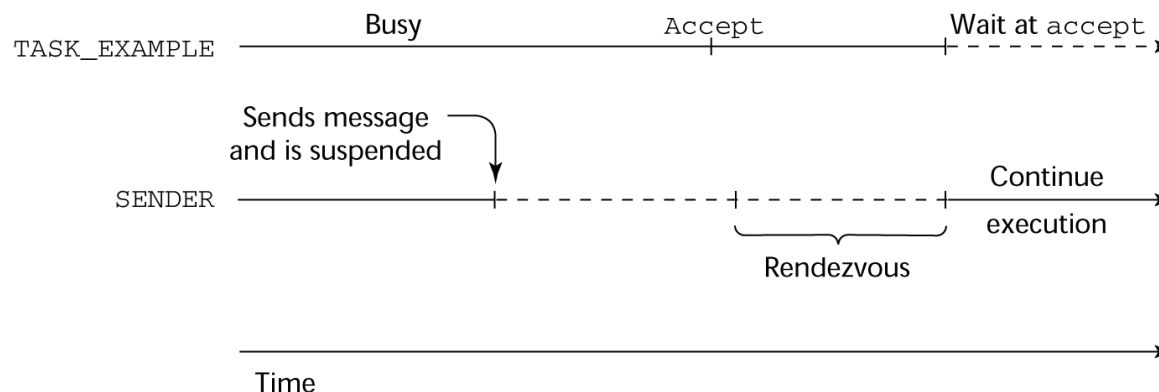


# RENDEZVOUS

(NOT AS ROMANTIC AS IT SOUNDS!)



(a) TASK\_EXAMPLE waits for SENDER



(b) SENDER waits for TASK\_EXAMPLE

# CPU SCHEDULING

**How does the OS decide which of several processes to take off the ready queue?**


**Scheduling: deciding which processes are given access to resources from moment to moment.**

# ASSUMPTION: CPU BURSTS

**Execution model: processes alternate between bursts of CPU and I/O**

- Process typically uses the CPU for some period of time, then does I/O, then uses CPU again
- Each scheduling decision is about which process to give to the CPU for use by its next CPU burst
- With time slicing, process may be forced to give up CPU before finishing current CPU burst.

# WHAT IS IMPORTANT IN A SCHEDULING ALGORITHM?



# WHAT IS IMPORTANT IN A SCHEDULING ALGORITHM?

## Minimize Response Time

- Elapsed time to do an operation (job)
- Response time is what the user sees
  - Time to echo keystroke in editor
  - Time to compile a program
  - Real-time Tasks: Must meet deadlines imposed by World

# WHAT IS IMPORTANT IN A SCHEDULING ALGORITHM?

## Maximize Throughput

- Jobs per second
- Throughput related to response time, but not identical
  - Minimizing response time will lead to more context switching than if you maximized only throughput
- Minimize overhead (context switch time) as well as efficient use of resources (CPU, disk, memory, etc.)

# WHAT IS IMPORTANT IN A SCHEDULING ALGORITHM?

## Fairness

- Share CPU among users or processes in some equitable way
- Not just minimizing average response time

# **SCHEDULING ALGORITHMS: FIRST-COME, FIRST-SERVED (FCFS)**

**“Run until Done:” FIFO algorithm**

**In the beginning, this meant one process runs non-preemptively until it is finished**

- Including any blocking for I/O operations

**Now, FCFS means that a process keeps the CPU until it completes its burst**

**Example: Three processes arrive in order P1, P2, P3.**

- P1 burst time: 24
- P2 burst time: 3
- P3 burst time: 3


**Draw the Gantt Chart and compute Average Waiting Time and Average Completion Time.**



# SCHEDULING ALGORITHMS: FIRST-COME, FIRST-SERVED (FCFS)

**Example: Three processes arrive in order P1, P2, P3.**

- P1 burst time: 24
- P2 burst time: 3
- P3 burst time: 3



**Waiting Time**

- P1: 0
- P2: 24
- P3: 27

**Completion Time:**

- P1: 24
- P2: 27
- P3: 30

**Average Waiting Time:  $(0+24+27)/3 = 17$**

**Average Completion Time:  $(24+27+30)/3 = 27$**

# SCHEDULING ALGORITHMS: FIRST-COME, FIRST-SERVED (FCFS)


**What if their order had been P2, P3, P1?**

- P1 burst time: 24
- P2 burst time: 3
- P3 burst time: 3

# SCHEDULING ALGORITHMS: FIRST-COME, FIRST-SERVED (FCFS)

What if their order had been P2, P3, P1?

- P1 burst time: 24
- P2 burst time: 3
- P3 burst time: 3



**Waiting Time**

- P2: 0
- P3: 3
- P1: 6

**Completion Time:**

- P2: 3
- P3: 6
- P1: 30

**Average Waiting Time:**  $(0+3+6)/3 = 3$  (compared to 17)

**Average Completion Time:**  $(3+6+30)/3 = 13$  (compared to 27)

# SCHEDULING ALGORITHMS: FIRST-COME, FIRST-SERVED (FCFS)


**Average Waiting Time:  $(0+3+6)/3 = 3$  (compared to 17)**

**Average Completion Time:  $(3+6+30)/3 = 13$  (compared to 27)**

## **FCFS Pros and Cons:**

- Simple (+)
- Short jobs get stuck behind long ones (-)
- Performance is highly dependent on the order in which jobs arrive (-)

# HOW CAN WE IMPROVE ON THIS?



# ROUND ROBIN (RR) SCHEDULING

## Round Robin Scheme

- Each process gets a small unit of CPU time (time quantum)
  - Usually 10-100 ms
- After quantum expires, the process is preempted and added to the end of the ready queue
- Suppose  $N$  processes in ready queue and time quantum is  $Q$  ms:
  - Each process gets  $1/N$  of the CPU time
  - In chunks of at most  $Q$  ms
  - What is the maximum wait time for each process?
    - ***No process waits more than  $(N-1)Q$  time units***

# ROUND ROBIN (RR) SCHEDULING

Performance Depends on the **value** of Q

- Small Q => interleaved
- Large Q is like **FCFS**
- Q must be large with respect to **context switch time**, otherwise overhead is too high
  - Spending most of your time context switching!

# EXAMPLE OF RR WITH TIME QUANTUM = 4


## Process Burst Time

$P_1$  24

$P_2$  3

$P_3$  3

The Gantt chart is:





# EXAMPLE OF RR WITH TIME QUANTUM = 4

## Process

## Burst Time

$P_1$

24

$P_2$

3

$P_3$


3

## Waiting Time:

- $P_1: (10-4) = 6$
- $P_2: (4-0) = 4$
- $P_3: (7-0) = 7$

## Completion Time:

- $P_1: 30$
- $P_2: 7$
- $P_3: 10$



Average Waiting Time:  $(6 + 4 + 7)/3 = 5.67$


Average Completion Time:  $(30+7+10)/3=15.67$

# EXAMPLE OF RR WITH TIME QUANTUM = 20

<u>Process</u>	<u>Burst Time</u>
<i>P1</i>	53
<i>P2</i>	8
<i>P3</i>	68
<i>P4</i>	24

# EXAMPLE OF RR WITH TIME QUANTUM = 20

<u>Process</u>	<u>Burst Time</u>
P1	<del>5</del> 3 <del>3</del> 3 <del>1</del> 3 0
P2	<del>8</del> 0
P3	<del>6</del> 8 <del>4</del> 8 <del>2</del> 8 <del>8</del> 0
P4	<del>2</del> 4 <del>4</del> 0




# EXAMPLE OF RR WITH TIME QUANTUM = 20

## Completion Time:

- P1: 125
- P2: 28
- P3: 153
- P4: 112

<u>Process</u>	<u>Burst Time</u>
<i>P1</i>	53
<i>P2</i>	8
<i>P3</i>	68
<i>P4</i>	24




**Average Completion Time:  $(125+28+153+112)/4 = 104.5$**

# EXAMPLE OF RR WITH TIME QUANTUM = 20

## Waiting Time:

- For P1:
  - $(68-20) + (112-88) = 72$

<u>Process</u>	<u>Burst Time</u>
P1	53
P2	8
P3	68
P4	24




# EXAMPLE OF RR WITH TIME QUANTUM = 20

## Waiting Time:

- For P1:
  - $(68-20)+(112-88) = 72$
- For P2:
  - $20-0 = 20$

<u>Process</u>	<u>Burst Time</u>
P1	53
P2	8
P3	68
P4	24




# EXAMPLE OF RR WITH TIME QUANTUM = 20

## Waiting Time:

- For P1:
  - $(68-20)+(112-88) = 72$
- For P2:
  - $20-0 = 20$
- For P3:
  - $(28-0)+(88-48)+(125-108) = 85$

<u>Process</u>	<u>Burst Time</u>
P1	53
P2	8
P3	68
P4	24



# EXAMPLE OF RR WITH TIME QUANTUM = 20

## Waiting Time:

- For P1:
  - $(68-20)+(112-88) = 72$
- For P2:
  - $20-0 = 20$
- For P3:
  - $(28-0)+(88-48)+(125-108) = 85$
- For P4:
  - $(48-0)+(108-68) = 88$

<u>Process</u>	<u>Burst Time</u>
P1	53
P2	8
P3	68
P4	24

P <sub>1</sub>	P <sub>2</sub>	P <sub>3</sub>	P <sub>4</sub>	P <sub>1</sub>	P <sub>3</sub>	P <sub>4</sub>	P <sub>1</sub>	P <sub>3</sub>	P <sub>3</sub>	
0	20	28	48	68	88	108	112	125	145	153

**Average Waiting Time:  $(72+20+85+88)/4 = 66.25$**



# RR SUMMARY

## Pros and Cons:

- Better for short jobs (+)
- Fair (+)
- Context-switching time adds up for long jobs (-)

## If the chosen quantum is

- too large, response time suffers
- infinite, performance is the same as FCFS
- too small, throughput suffers as percentage of overhead grows

# **THANK YOU!**

## **QUESTIONS?**